

# Bigdata and Hadoop

Haridas Narayanaswamy

<https://haridas.in>

06/Mar/2019

# Quick reminder

1. HYD and BLR: Online Attendance registration link

<https://docs.google.com/spreadsheets/d/1DAuclSqS6xdv4QtixUQIs6OjpxXRLVEEmnCeAkgD9z8/edit#gid=0>

2.

```
docker pull haridasn/hadoop-2.8.5
```

```
docker pull haridasn/hadoop-cli
```



# Agenda

- Introduction to the big-data problems
- How we can scale systems.
- Hadoop introduction
- Setup a Hadoop cluster on your laptop.



# IO bound vs CPU bound problems

- Downloading a file
- Bitcoin mining ( proof of work )
- Watching a movie
- ETL jobs
- ML training

# Latency Numbers

- nano second - unit speed

## Latency Comparison Numbers (~2012)

L1 cache reference	0.5	ns			
Branch mispredict	5	ns			
L2 cache reference	7	ns			14x L1 cache
Mutex lock/unlock	25	ns			
Main memory reference	100	ns			20x L2 cache, 200x L1 cache
Compress 1K bytes with Zippy	3,000	ns	3	us	
Send 1K bytes over 1 Gbps network	10,000	ns	10	us	
Read 4K randomly from SSD*	150,000	ns	150	us	~1GB/sec SSD
Read 1 MB sequentially from memory	250,000	ns	250	us	
Round trip within same datacenter	500,000	ns	500	us	
Read 1 MB sequentially from SSD*	1,000,000	ns	1,000	us	1 ms ~1GB/sec SSD, 4X memory
Disk seek	10,000,000	ns	10,000	us	10 ms 20x datacenter roundtrip
Read 1 MB sequentially from disk	20,000,000	ns	20,000	us	20 ms 80x memory, 20X SSD
Send packet CA→Netherlands→CA	150,000,000	ns	150,000	us	150 ms



# Application Architectures

Online  
Standalone  
Distributed


# Monolithic

- All on single machine
- Every application starts here



# N-tier systems


- Multiple services in your application
- Separation of concerns
- Master-slave database systems
- Load balancer
- Caching layer
- Pretty good for standard application workloads.





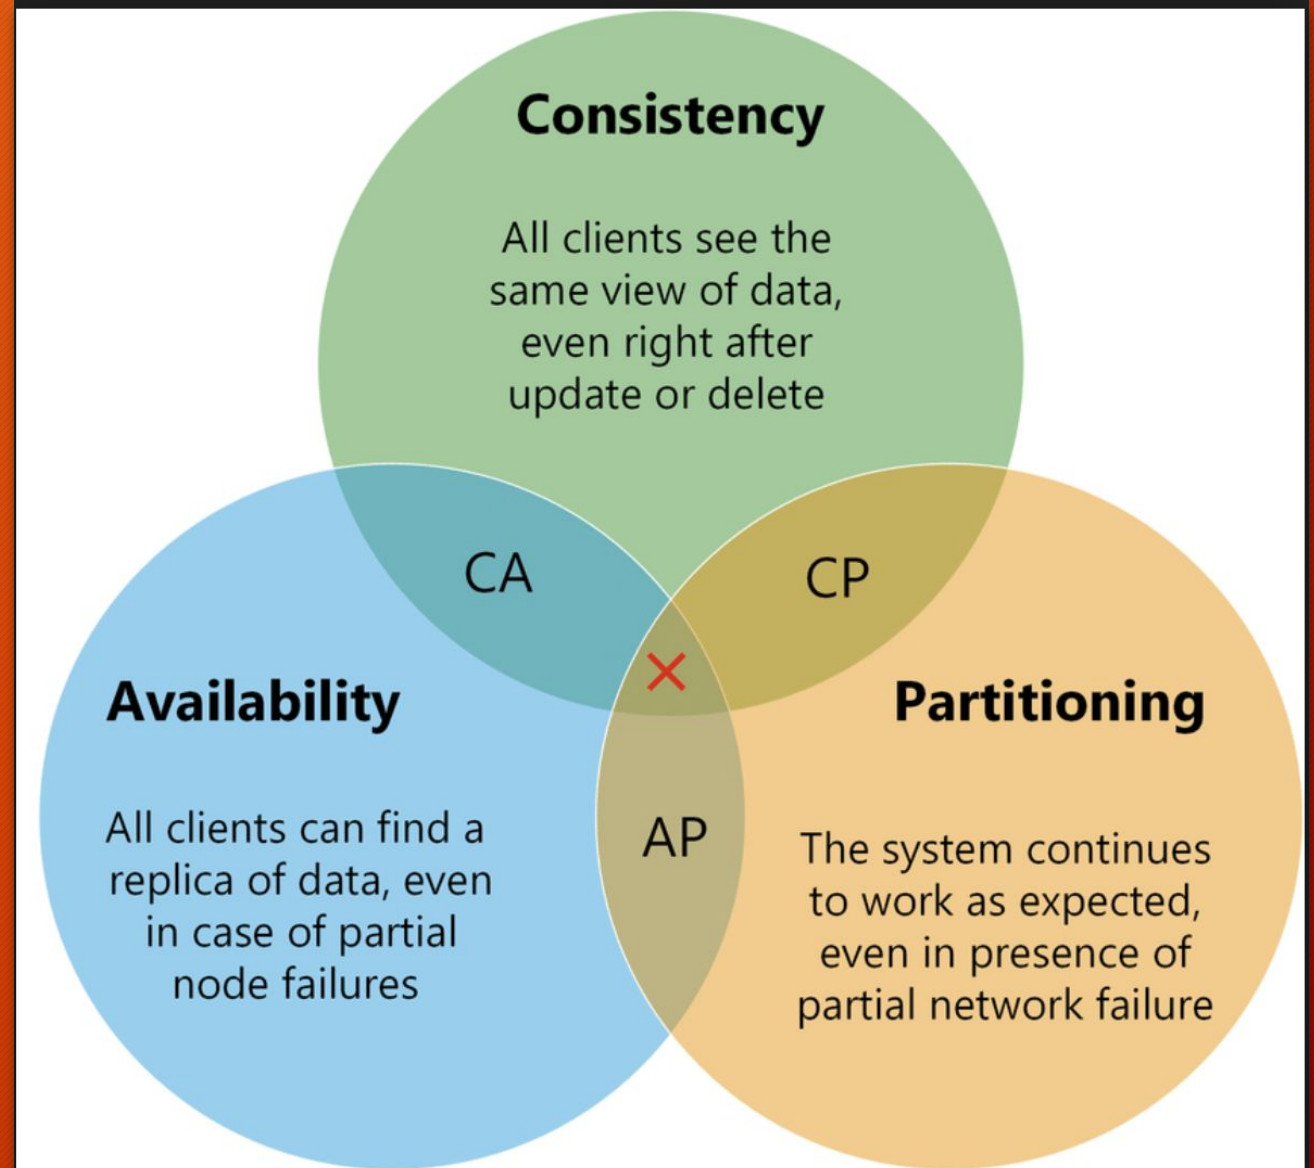
# SOA and Microservices

- Service Oriented Architecture
- Each service will do only one task - fine grained separation, hence named as micro-services
- API Gateways
- Messaging Systems
- Scalable Databases accessible to micro services
- Caching services
- Etc.



# CAP Theorem

- **Consistency** - Reads from all machines on the cluster would give same data.
- **Availability** - All operation would produce some result, even though they aren't consistent.
- **Partition tolerance** - System perform normally even after some nodes got disconnected from network.
- Eventual Consistent systems
- Split brain problem





# Bigdata platforms

Where they fit in ?  
Offline vs Online systems



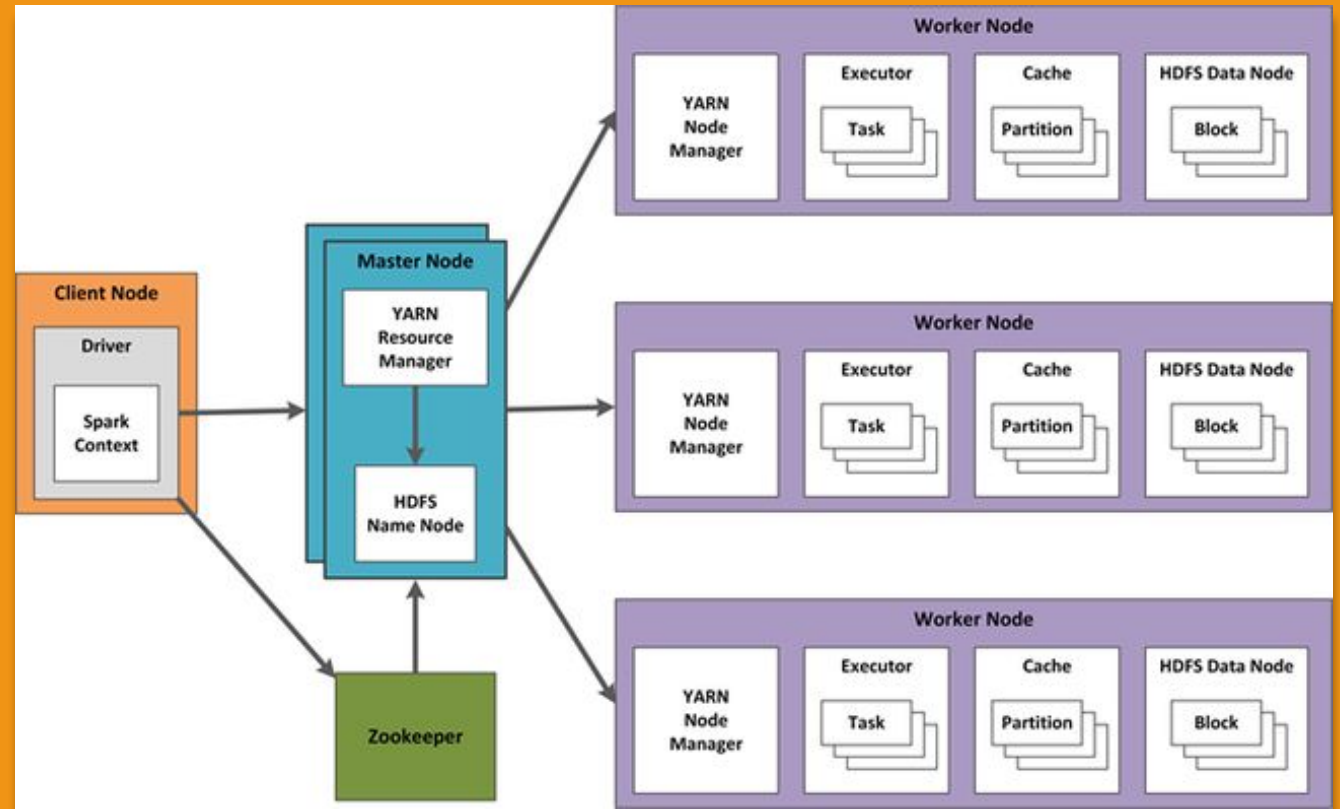
# Count number of unique hits from India

- Consider this scenario for Facebook
- Daily total log file comes is in Petabytes
- One machine can't save it
- Other scenarios
  - User click tracking
  - Tracking effectiveness of an Ad
- All flavours of ETL jobs



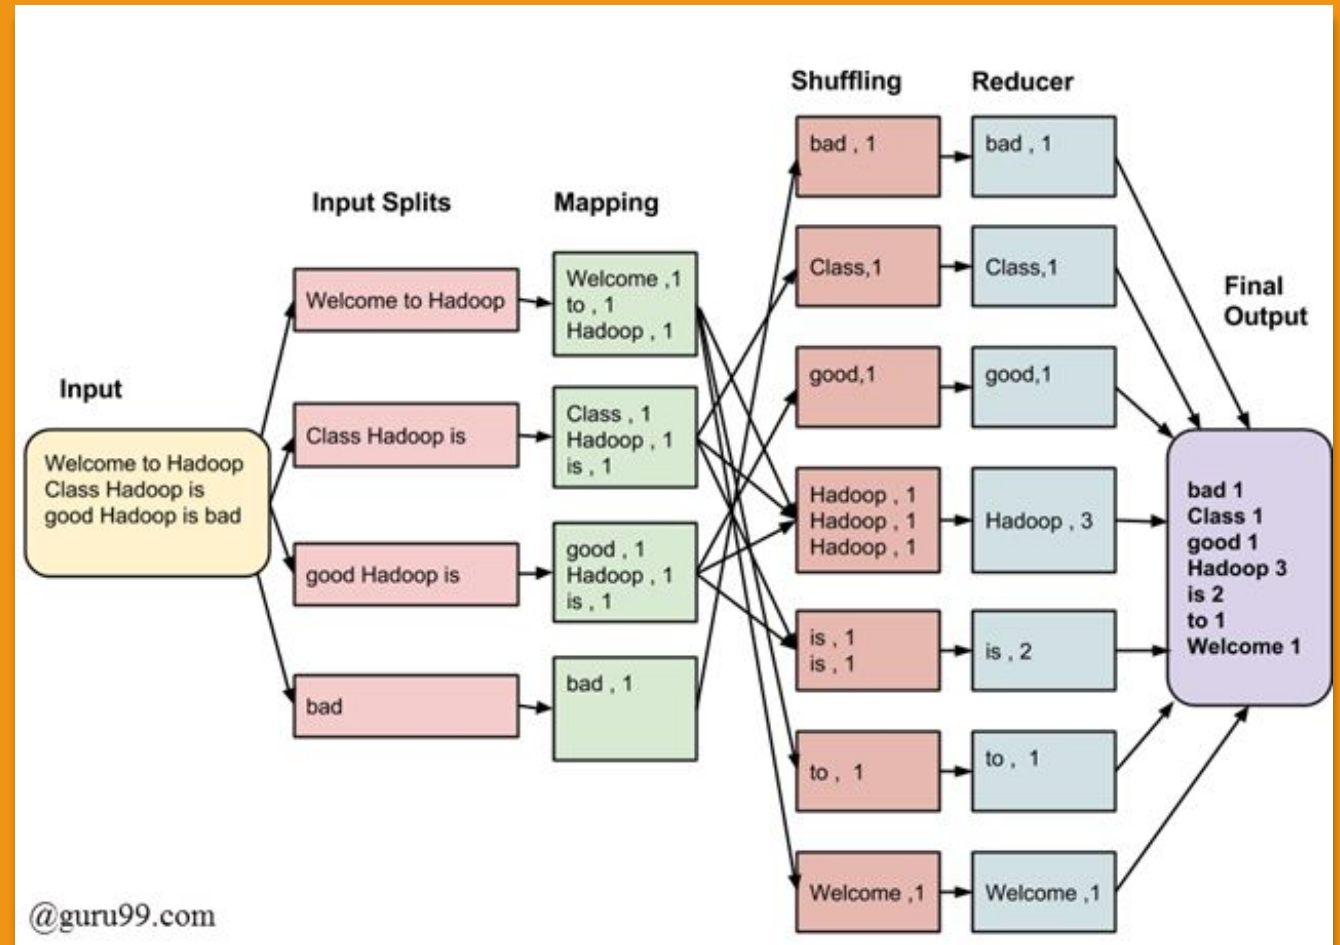
# Hadoop

- Every node can store the data - HDFS
- Every node can do processing on data it holds locally.
- Easily scalable
- Redundant and fault tolerant
- Main Components
  - Name Node
  - Data Node
  - Resource Manager
  - Node Manager



# Map-Reduce compute model

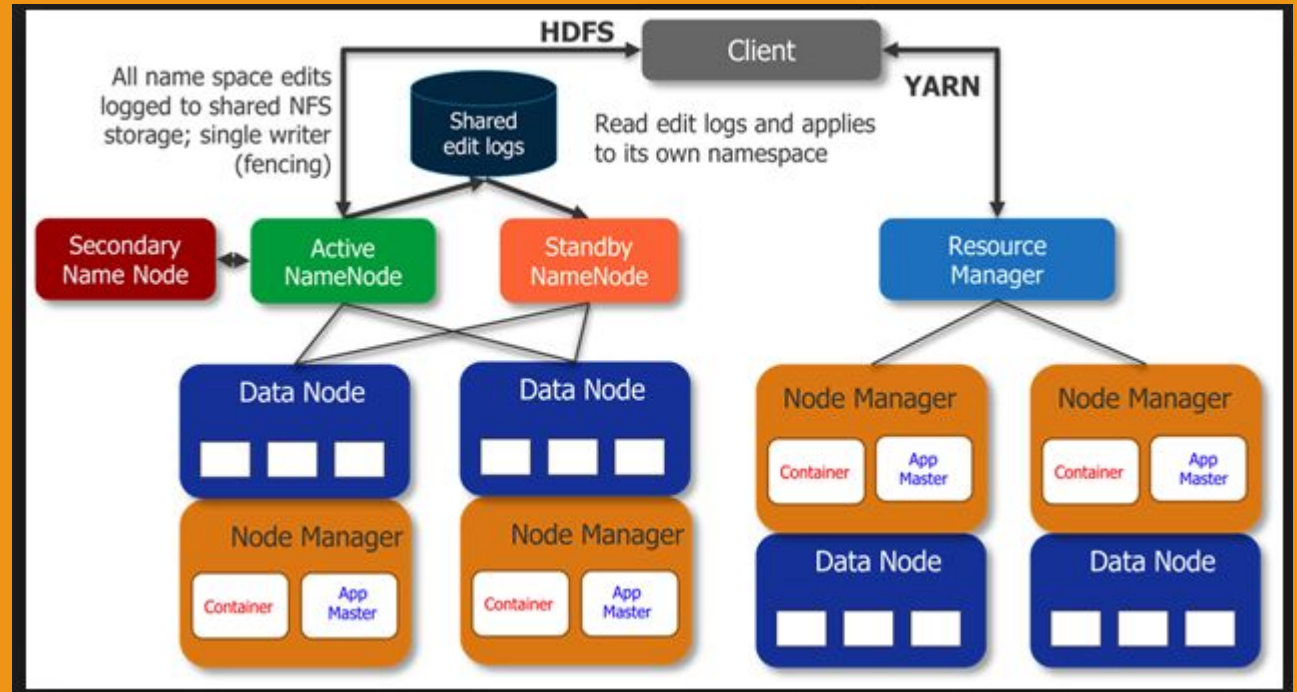
- Main stages of map-reduce
- Copy code not data.
- Data locality
- How we can use map-reduce to count unique hits from a region ?





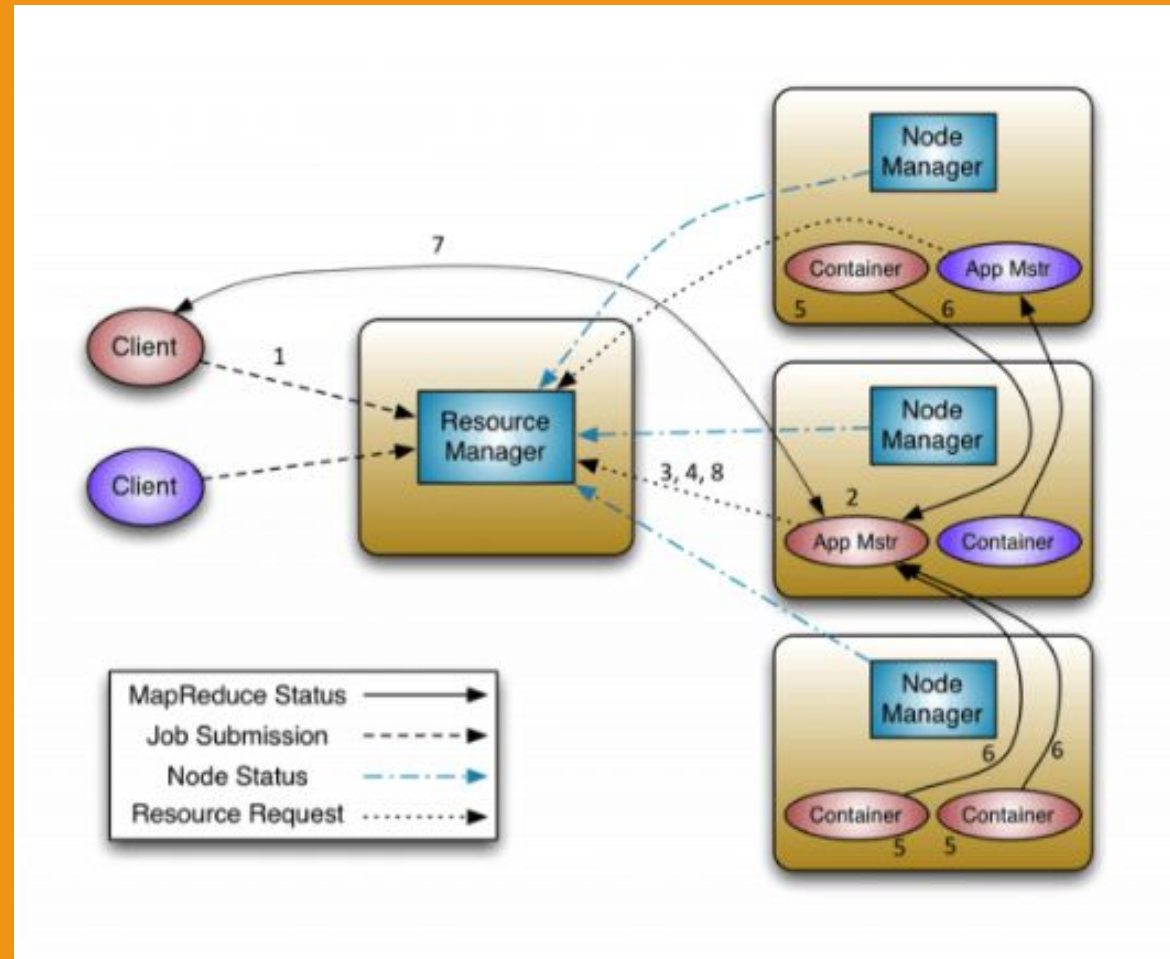
# Yarn and HDFS

- Main stages of map-reduce
- Dynamic programming ?
- Bring code to data location
- Data locality
- How we can use mapreduce to count unique hits from a region ?



# Yarn Framework

- Client Submits jobs into cluster
- AM handles application orchestration, once it has been started by RM
- Containers are actual resources (CPU/Ram/IO) allocated to AM.
- Job progress tracking





# Hybrid environments

- Storage on cloud storages s3/azure storage
- Execution engine on our cloud or Kubernetes.

Workshop



# Next: Spark on Hadoop Yarn and Pyspark

ON 13-March-2019

# Thank you

Haridas N <hn@haridas.in>